

```

In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

In [2]: import tensorflow as tf
from tensorflow import keras

In [3]: from sklearn.datasets import fetch_california_housing

In [4]: df = fetch_california_housing()

In [5]: print(df.feature_names)

['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude']

In [6]: from sklearn.model_selection import train_test_split

In [7]: x_train_full,x_test,y_train_full,y_test =train_test_split(df.data,df.target,random_state = 42)
x_train,x_valid,y_train,y_valid = train_test_split(x_train_full,y_train_full,random_state=42)

In [8]: from sklearn.preprocessing import StandardScaler

In [9]: scaler = StandardScaler()

In [10]: x_train = scaler.fit_transform(x_train)
x_valid = scaler.transform(x_valid)
x_test = scaler.transform(x_test)

In [11]: np.random.seed(42)
tf.random.set_seed(42)

In [12]: x_train.shape

Out[12]: (11610, 8)

In [13]: model = keras.models.Sequential([
    keras.layers.Dense(30,activation = 'relu',input_shape = [8]),
    keras.layers.Dense(30,activation = 'relu'),
    keras.layers.Dense(1)
])

In [14]: model.summary()

Model: "sequential"

Layer (type)                Output Shape                Param #
=====
dense (Dense)                (None, 30)                  270
dense_1 (Dense)              (None, 30)                  930
dense_2 (Dense)              (None, 1)                   31
=====
Total params: 1,231
Trainable params: 1,231
Non-trainable params: 0

In [15]: model.compile(loss = 'mean_squared_error',
    optimizer = keras.optimizers.SGD(),
    metrics = ['mae'])

In [16]: model_history = model.fit(x_train,y_train,epochs = 30,validation_data = (x_valid,y_valid))

Epoch 1/30
363/363 [=====] - 3s 4ms/step - loss: 0.6860 - mae: 0.5876 - val_loss: 0.6226 - val_mae: 0.4769
Epoch 2/30
363/363 [=====] - 1s 3ms/step - loss: 0.4286 - mae: 0.4700 - val_loss: 3.4166 - val_mae: 0.4771
Epoch 3/30
363/363 [=====] - 1s 3ms/step - loss: 0.4907 - mae: 0.4801 - val_loss: 2.1055 - val_mae: 0.4678
Epoch 4/30
363/363 [=====] - 1s 3ms/step - loss: 0.3964 - mae: 0.4451 - val_loss: 0.4980 - val_mae: 0.4321
Epoch 5/30
363/363 [=====] - 1s 3ms/step - loss: 0.3761 - mae: 0.4329 - val_loss: 0.3935 - val_mae: 0.4251
Epoch 6/30
363/363 [=====] - 1s 3ms/step - loss: 0.3636 - mae: 0.4265 - val_loss: 0.3986 - val_mae:

```

```

e: 0.4172
Epoch 7/30
363/363 [=====] - 1s 3ms/step - loss: 0.3592 - mae: 0.4226 - val_loss: 0.3733 - val_mae: 0.4131
Epoch 8/30
363/363 [=====] - 1s 3ms/step - loss: 0.3504 - mae: 0.4175 - val_loss: 0.3645 - val_mae: 0.4113
Epoch 9/30
363/363 [=====] - 1s 3ms/step - loss: 0.3447 - mae: 0.4131 - val_loss: 0.3582 - val_mae: 0.3986
Epoch 10/30
363/363 [=====] - 1s 3ms/step - loss: 0.3416 - mae: 0.4117 - val_loss: 0.3622 - val_mae: 0.4074
Epoch 11/30
363/363 [=====] - 1s 3ms/step - loss: 0.3389 - mae: 0.4088 - val_loss: 0.3334 - val_mae: 0.3937
Epoch 12/30
363/363 [=====] - 1s 3ms/step - loss: 0.3371 - mae: 0.4072 - val_loss: 0.3537 - val_mae: 0.3943
Epoch 13/30
363/363 [=====] - 1s 3ms/step - loss: 0.3325 - mae: 0.4049 - val_loss: 0.3405 - val_mae: 0.3889
Epoch 14/30
363/363 [=====] - 1s 3ms/step - loss: 0.3280 - mae: 0.4017 - val_loss: 0.3269 - val_mae: 0.3876
Epoch 15/30
363/363 [=====] - 1s 3ms/step - loss: 0.3293 - mae: 0.4013 - val_loss: 0.3321 - val_mae: 0.3928
Epoch 16/30
363/363 [=====] - 1s 3ms/step - loss: 0.3234 - mae: 0.3994 - val_loss: 0.3235 - val_mae: 0.3912
Epoch 17/30
363/363 [=====] - 1s 3ms/step - loss: 0.3215 - mae: 0.3976 - val_loss: 0.3576 - val_mae: 0.4040
Epoch 18/30
363/363 [=====] - 1s 3ms/step - loss: 0.3210 - mae: 0.3962 - val_loss: 0.3118 - val_mae: 0.3799
Epoch 19/30
363/363 [=====] - 1s 3ms/step - loss: 0.3224 - mae: 0.3956 - val_loss: 0.3251 - val_mae: 0.3821
Epoch 20/30
363/363 [=====] - 1s 3ms/step - loss: 0.3171 - mae: 0.3931 - val_loss: 0.3686 - val_mae: 0.4010
Epoch 21/30
363/363 [=====] - 1s 3ms/step - loss: 0.3137 - mae: 0.3914 - val_loss: 0.3139 - val_mae: 0.3946
Epoch 22/30
363/363 [=====] - 1s 3ms/step - loss: 0.3120 - mae: 0.3906 - val_loss: 0.3211 - val_mae: 0.3843
Epoch 23/30
363/363 [=====] - 1s 3ms/step - loss: 0.3104 - mae: 0.3883 - val_loss: 0.3098 - val_mae: 0.3751
Epoch 24/30
363/363 [=====] - 1s 3ms/step - loss: 0.3101 - mae: 0.3891 - val_loss: 0.3160 - val_mae: 0.3832
Epoch 25/30
363/363 [=====] - 1s 3ms/step - loss: 0.3078 - mae: 0.3866 - val_loss: 0.3335 - val_mae: 0.3765
Epoch 26/30
363/363 [=====] - 1s 3ms/step - loss: 0.3066 - mae: 0.3858 - val_loss: 0.3140 - val_mae: 0.3801
Epoch 27/30
363/363 [=====] - 1s 3ms/step - loss: 0.3063 - mae: 0.3854 - val_loss: 0.3009 - val_mae: 0.3741
Epoch 28/30
363/363 [=====] - 1s 3ms/step - loss: 0.3062 - mae: 0.3838 - val_loss: 0.3166 - val_mae: 0.3818
Epoch 29/30
363/363 [=====] - 1s 3ms/step - loss: 0.3081 - mae: 0.3855 - val_loss: 0.3078 - val_mae: 0.3860
Epoch 30/30
363/363 [=====] - 1s 3ms/step - loss: 0.3069 - mae: 0.3845 - val_loss: 0.3074 - val_mae: 0.3743

```

```
In [17]: mae_test = model.evaluate(x_test,y_test)
```

```
162/162 [=====] - 0s 2ms/step - loss: 0.3095 - mae: 0.3819
```

```
In [18]: model_history.history
```


```
Out[18]: {'loss': [0.6860350966453552,
0.42857858538627625,
0.49068236351013184,
0.39637333154678345,
0.37606489658355713,
0.36363890767097473,
0.3591632544994354,
0.3504336178302765,
0.3447352349758148,
```

0.3416411280632019,
0.3388660252094269,
0.33707621693611145,
0.3325381278991699,
0.32801133394241333,
0.32931452989578247,
0.3233952522277832,
0.32146942615509033,
0.3209802508354187,
0.32241880893707275,
0.3171079456806183,
0.3137245774269104,
0.3120401203632355,
0.3104282021522522,
0.3100728392601013,
0.30784472823143005,
0.30662301182746887,
0.30626380443573,
0.3062289357185364,
0.30806034803390503,
0.30694344639778137],
'mae': [0.5875542759895325,
0.4699855148792267,
0.48008817434310913,
0.44509610533714294,
0.4329044818878174,
0.426518052816391,
0.4226357340812683,
0.41745978593826294,
0.4131312370300293,
0.41174769401550293,
0.40878793597221375,
0.4072113037109375,
0.40491777658462524,
0.40173640847206116,
0.40128058195114136,
0.39936044812202454,
0.39757245779037476,
0.39616405963897705,
0.3955683708190918,
0.39313799142837524,
0.3913708031177521,
0.3906291127204895,
0.38829725980758667,
0.38911840319633484,
0.38659846782684326,
0.38575366139411926,
0.3853800892829895,
0.38375401496887207,
0.3855398893356323,
0.3844830095767975],
'val_loss': [0.6225500702857971,
3.416584014892578,
2.105531930923462,
0.4979768991470337,
0.39347290992736816,
0.3986354172229767,
0.3732934296131134,
0.3645224869251251,
0.35822364687919617,
0.3621854782104492,
0.33341464400291443,
0.35373955965042114,
0.3404729962348938,
0.32689765095710754,
0.3320634663105011,
0.3235388696193695,
0.3575971722602844,
0.3117719292640686,
0.32507187128067017,
0.36857613921165466,
0.31391409039497375,
0.32108721137046814,
0.30976879596710205,
0.31599611043930054,
0.3334956765174866,
0.3139708936214447,
0.30089184641838074,
0.3165510594844818,
0.3077872395515442,
0.3073861300945282],
'val_mae': [0.47691401839256287,
0.4771106243133545,
0.46782588958740234,
0.43206143379211426,
0.4250621199607849,
0.4172133207321167,
0.413072794675827,
0.4112563729286194,

```
0.39858952164649963,
0.40735992789268494,
0.39371392130851746,
0.394314706325531,
0.38892585039138794,
0.3875541388988495,
0.39283493161201477,
0.3912133276462555,
0.403969407081604,
0.37989693880081177,
0.38205376267433167,
0.4010195732116699,
0.3946186304092407,
0.38433516025543213,
0.3751460611820221,
0.3832392990589142,
0.37651559710502625,
0.3801248073577881,
0.3740726113319397,
0.3817773759365082,
0.38596710562705994,
0.3742929697036743]}
```

```
In [19]: pd.DataFrame(model_history.history).plot(figsize = (8,5))
plt.grid(True)

plt.show()
```



```
In [20]: x_new = x_test[:3]
```

```
In [21]: y_pred = model.predict(x_new)

1/1 [=====] - 0s 236ms/step
```

```
In [22]: y_pred
```

```
Out[22]: array([[0.6329897],
                [1.3350961],
                [4.173784 ]], dtype=float32)
```

```
In [23]: y_test[:3]
```

```
Out[23]: array([0.477 , 0.458 , 5.00001])
```

Functional API

Not all neural network models are simply sequential. Some may have complex topologies. Some may have multiple inputs and/or multiple outputs. For example a wide and deep learning connections all or path of inputs directly to the output layer

```
In [24]: del model
```

```
In [25]: keras.backend.clear_session()
```

```
In [29]: input_ = keras.layers.Input(shape = x_train.shape[1:])
hidden1 = keras.layers.Dense(30, activation = 'relu')(input_)
hidden2 = keras.layers.Dense(30, activation = 'relu')(hidden1)
concat = keras.layers.concatenate([input_, hidden2])
output = keras.layers.Dense(1)(concat)
```

```
model = keras.models.Model(inputs = [input_],outputs = [output])
```

```
In [30]: model.summary()
```

Model: "model"

| Layer (type) | Output Shape | Param # | Connected to |
|-----------------------------|--------------|---------|---------------------------------------|
| input_4 (InputLayer) | (None, 8) | 0 | [] |
| dense_7 (Dense) | (None, 30) | 270 | ['input_4[0][0]'] |
| dense_8 (Dense) | (None, 30) | 930 | ['dense_7[0][0]'] |
| concatenate_2 (Concatenate) | (None, 38) | 0 | ['input_4[0][0]', 'dense_8[0][0]'] |
| dense_9 (Dense) | (None, 1) | 39 | ['concatenate_2[0][0]'] |

=====
Total params: 1,239
Trainable params: 1,239
Non-trainable params: 0
=====

```
In [ ]:
```